

Import libraries

```
In [1]: import numpy as np
import pandas as pd
from math import sqrt
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
from sklearn.model_selection import GridSearchCV
```

Load the dataset

```
In [2]: df = pd.read_csv("concrete.csv")
df.head(10)
```

Out[2]:

	Cement (component 1)(kg in a m ³ mixture)	Blast Furnace Slag (component 2)(kg in a m ³ mixture)	Fly Ash (component 3)(kg in a m ³ mixture)	Water (component 4)(kg in a m ³ mixture)	Superplasticizer (component 5) (kg in a m ³ mixture)	Coarse Aggregate (component 6)(kg in a m ³ mixture)	Fine Aggregate (component 7)(kg in a m ³ mixture)
0	540.0	0.0	0.0	162.0	2.5	1040.0	676.0
1	540.0	0.0	0.0	162.0	2.5	1055.0	676.0
2	332.5	142.5	0.0	228.0	0.0	932.0	594.0
3	332.5	142.5	0.0	228.0	0.0	932.0	594.0
4	198.6	132.4	0.0	192.0	0.0	978.4	825.5
5	266.0	114.0	0.0	228.0	0.0	932.0	670.0
6	380.0	95.0	0.0	228.0	0.0	932.0	594.0
7	380.0	95.0	0.0	228.0	0.0	932.0	594.0
8	266.0	114.0	0.0	228.0	0.0	932.0	670.0
9	475.0	0.0	0.0	228.0	0.0	932.0	594.0

Data Preprocessing

```
In [3]: df1 = df.copy()
```

```
In [4]: df1.isnull().sum()
```

```
Out[4]: Cement (component 1)(kg in a m^3 mixture)      0
Blast Furnace Slag (component 2)(kg in a m^3 mixture)  0
Fly Ash (component 3)(kg in a m^3 mixture)            0
Water (component 4)(kg in a m^3 mixture)              0
Superplasticizer (component 5)(kg in a m^3 mixture)   0
Coarse Aggregate (component 6)(kg in a m^3 mixture)   0
Fine Aggregate (component 7)(kg in a m^3 mixture)     0
Age (day)                                              0
strength                                              0
dtype: int64
```

```
In [5]: df1.describe()
```

```
Out[5]:
```

	Cement (component 1)(kg in a m^3 mixture)	Blast Furnace Slag (component 2)(kg in a m^3 mixture)	Fly Ash (component 3)(kg in a m^3 mixture)	Water (component 4)(kg in a m^3 mixture)	Superplasticizer (component 5) (kg in a m^3 mixture)	Coarse Aggregate (component 6)(kg in a m^3 mixture)	Aggr (comp 7)(k mi
count	1030.000000	1030.000000	1030.000000	1030.000000	1030.000000	1030.000000	1030.000000
mean	281.167864	73.895825	54.188350	181.567282	6.204660	972.918932	773.500000
std	104.506364	86.279342	63.997004	21.354219	5.973841	77.753954	80.100000
min	102.000000	0.000000	0.000000	121.800000	0.000000	801.000000	594.000000
25%	192.375000	0.000000	0.000000	164.900000	0.000000	932.000000	730.900000
50%	272.900000	22.000000	0.000000	185.000000	6.400000	968.000000	779.500000
75%	350.000000	142.950000	118.300000	192.000000	10.200000	1029.400000	824.000000
max	540.000000	359.400000	200.100000	247.000000	32.200000	1145.000000	992.600000

```
In [6]: x = df1.drop('strength', axis=1)
y = df1['strength']
```

```
In [7]: x.shape
```

```
Out[7]: (1030, 8)
```

```
In [8]: y.shape
```

```
Out[8]: (1030,)
```

```
In [9]: x1 = x.copy()
```

Applied Normalization

```
In [10]: m = MinMaxScaler()
```

```
In [11]: for col in x1.columns:
          if x1[col].dtype == object:
              continue
          else:
              m.fit(x1[[col]])
              x1[col] = m.transform(x1[[col]])
          x1.head()
```

Out[11]:

	Cement (component 1)(kg in a m ³ mixture)	Blast Furnace Slag (component 2)(kg in a m ³ mixture)	Fly Ash (component 3)(kg in a m ³ mixture)	Water (component 4)(kg in a m ³ mixture)	Superplasticizer (component 5) (kg in a m ³ mixture)	Coarse Aggregate (component 6)(kg in a m ³ mixture)	Fine Aggregate (component 7)(kg in a m ³ mixture)
0	1.000000	0.000000	0.0	0.321086	0.07764	0.694767	0.205720
1	1.000000	0.000000	0.0	0.321086	0.07764	0.738372	0.205720
2	0.526256	0.396494	0.0	0.848243	0.00000	0.380814	0.000000
3	0.526256	0.396494	0.0	0.848243	0.00000	0.380814	0.000000
4	0.220548	0.368392	0.0	0.560703	0.00000	0.515698	0.580783

```
In [12]: xtrain, xtest, ytrain, ytest = train_test_split(x1, y, test_size = 0.25, random_state=42)
```

Linear Regression

```
In [13]: lr = LinearRegression()
          lr.fit(xtrain, ytrain)
```

Out[13]: LinearRegression()

```
In [14]: ypred1 = lr.predict(xtest)
ypred1
```

```
Out[14]: array([59.09859859, 51.97626903, 63.39987251, 51.51583827, 17.12396228,
39.46825434, 26.47328415, 44.77707843, 29.63067645, 37.93010481,
27.84279166, 19.54919141, 66.81785117, 52.21871581, 29.94676466,
44.21751379, 29.05701223, 26.44812644, 31.8806001 , 32.08771228,
36.75906521, 31.75134132, 38.19736284, 25.03844605, 32.90412121,
34.07945783, 14.68126718, 40.1256464 , 41.83896327, 21.28573033,
35.7901225 , 30.79342806, 43.52433084, 45.50471383, 30.83600058,
29.33637776, 29.14495057, 38.51494716, 20.28393646, 38.56343404,
21.44355934, 15.88160473, 31.0992574 , 50.83828591, 20.63781584,
57.51048796, 50.65118732, 60.43076193, 20.18038732, 19.23597556,
40.20407262, 35.99947786, 29.66275818, 33.40748214, 46.74268335,
51.40978041, 28.0191688 , 16.01632264, 29.95463362, 18.43777596,
38.28982944, 20.06239506, 31.80326173, 55.46866206, 22.9228481 ,
21.31872757, 32.0692182 , 16.74554006, 25.70845259, 25.86659364,
17.84136298, 18.5558084 , 13.24490758, 27.79361218, 28.31322507,
20.36040383, 59.48918407, 44.5763007 , 54.25611871, 24.03059833,
43.94422864, 48.68078091, 32.08468576, 28.19353602, 68.12766917,
51.51583827, 34.32657578, 33.15241144, 26.88576597, 21.90400042,
26.92640271, 60.73133632, 24.13660383, 41.95617553, 39.25599808,
47.38219711, 28.05659423, 28.47193333, 27.08522273, 28.77614728,
30.44792376, 39.37763425, 45.85628288, 28.99251887, 76.94732786,
14.58624568, 41.21903421, 29.48096259, 33.55120961, 49.83028869,
64.66361278, 39.24230717, 26.91911628, 34.1838721 , 46.90017783,
55.687012 , 22.36673096, 40.61179234, 42.67863961, 39.28241845,
34.34180373, 23.25927658, 37.61285772, 28.00700775, 25.49486674,
31.72028893, 48.54192393, 54.98670093, 36.09690605, 31.99566875,
15.90222291, 34.4077842 , 14.80883001, 60.16488512, 14.6128627 ,
57.19243865, 26.300809 , 33.29516857, 32.3761724 , 33.46114948,
29.22487563, 32.86206925, 28.68058007, 26.65201491, 46.64182393,
36.5532213 , 27.79935504, 18.69447764, 19.82350571, 19.80866329,
36.01929186, 25.1139006 , 47.50880846, 27.61336797, 38.25604209,
33.96567437, 27.61475837, 67.64980041, 51.97626903, 40.13112602,
21.54531916, 58.44134821, 33.36283987, 49.33512452, 54.52627017,
52.54045462, 37.44951822, 25.66387898, 50.04673521, 33.61163254,
36.30651396, 25.69639372, 23.43834014, 34.19899734, 39.38102087,
20.52067401, 23.31944842, 39.51156236, 40.35395128, 41.6597914 ,
35.70977451, 26.2572127 , 49.72483738, 42.28036127, 37.21696018,
48.55395356, 21.50281227, 31.39251269, 46.65725394, 21.28244787,
47.0483777 , 37.51262928, 35.88190358, 47.74147428, 54.91389123,
31.3197729 , 40.43997549, 25.23181983, 22.39790058, 21.96645289,
38.22573415, 55.687012 , 17.49055494, 49.02457873, 53.38175107,
31.30009933, 34.1045647 , 48.88259533, 33.42882172, 31.93212094,
27.0905962 , 26.91573935, 25.07658529, 35.47135959, 30.05952485,
12.88480229, 19.26601094, 30.49840194, 23.17248479, 54.39353052,
48.55824528, 50.43181437, 49.31743931, 17.04326349, 40.09387719,
37.12023271, 31.64727865, 23.32079448, 25.98264227, 30.63032292,
22.09282847, 27.9995695 , 31.15366868, 36.08142838, 36.11580519,
39.13583935, 48.22035015, 28.38381183, 34.78128762, 25.61958138,
27.40734813, 37.51124279, 70.05740587, 20.42362276, 65.13770799,
35.03041566, 58.69362466, 16.34306056, 48.94453904, 60.62697772,
32.61069385, 34.38984241, 35.96331224, 26.73161831, 33.98249411,
26.26725833, 28.8628011 , 31.10480187]])
```

```
In [15]: lr.score(xtrain, ytrain)
```

```
Out[15]: 0.6099072868226489
```

```
In [16]: lr.score(xtest, ytest)
```

```
Out[16]: 0.6249829353885574
```

Random Forest

```
In [17]: rf = RandomForestRegressor()  
         rf.fit(xtrain, ytrain)
```

```
Out[17]: RandomForestRegressor()
```

```
In [18]: ypred2 = rf.predict(xtest)
ypred2
```

```
Out[18]: array([51.7677, 39.752, 72.0797, 34.7157, 11.6844,
44.3465, 24.593, 47.4667, 35.8479, 42.0349,
40.9161, 16.5403, 39.242, 36.6874, 24.4288,
23.0635, 38.2432, 18.1692, 38.3774, 31.4374,
36.2854, 36.2754, 45.7129, 10.7772, 34.5202,
37.9444, 11.3854, 45.0399, 53.7431, 14.611,
61.2266, 34.1811, 41.3641, 46.1804, 18.6616,
39.5868, 35.2933, 44.5758, 9.6715, 51.1208,
16.047, 6.2635, 40.0659, 48.7507, 12.7546,
65.3616, 52.86565833, 33.5199, 26.6673, 8.9335,
55.5573, 44.4205, 26.6544, 17.8486, 45.9052,
34.9508, 27.0334, 12.0979, 35.776, 19.9446,
44.7102, 14.0968, 35.7127, 51.06783333, 32.1173,
27.4283, 35.9702, 13.7235, 31.0047, 23.5758,
12.2764, 28.3087, 8.9397, 40.9209, 27.1012,
11.1705, 51.5817, 50.5152, 56.9296, 9.5529,
38.9344, 46.3576, 37.7475, 39.1114, 41.9161,
34.7157, 37.7354, 35.1399, 24.822, 22.3273,
32.1931, 71.4438, 13.1638, 54.9732, 39.2322,
47.5343, 22.1825, 37.6523, 20.5278, 32.4715,
31.45343333, 42.5982, 36.1215, 23.7369, 68.5768,
11.9361, 53.1343, 30.9063, 40.1442, 61.3533,
39.3766, 48.5578, 25.3408, 40.561, 34.68257714,
55.39967833, 19.6726, 35.0436, 57.1039, 39.3409,
19.6995, 27.1706, 52.7819, 30.6899, 27.9508,
42.925, 49.7774, 48.91333333, 43.8576, 35.1016,
12.0263, 35.97175714, 20.1487, 68.515, 10.1106,
41.2747, 32.3662, 39.4693, 21.62233333, 33.535,
28.4731, 33.2586, 15.4027, 32.5254, 43.8883,
42.1721, 28.3496, 14.8635, 7.0255, 17.8133,
42.1935, 24.7443, 41.36418, 22.0913, 42.6359,
40.376, 27.4221, 69.8712, 39.752, 47.2461,
13.7145, 59.0902, 35.3025, 52.0234, 33.9421,
65.2399, 41.7173, 22.9677, 39.24620667, 21.1953,
44.8119, 11.2205, 23.9534, 37.6054, 34.7409,
11.9858, 31.8325, 46.7167, 47.4168, 37.0238,
31.4922, 25.5559, 61.3893, 33.8828, 38.3854,
50.5768, 22.9817, 35.4308, 29.4718, 15.5033,
25.83907778, 38.6079, 39.7924, 38.86063333, 28.2652,
27.7279, 43.1227, 24.984, 23.3028, 28.5843,
39.8227, 55.39967833, 24.7257, 34.9554, 64.9485,
35.2817, 30.5917, 59.6019, 21.84876667, 35.7167,
15.5671, 30.7997, 20.1599, 44.5826, 42.2718,
5.3413, 13.5558, 37.5235, 23.7018, 71.2853,
29.24635, 54.9937, 59.005, 18.0472, 44.8915,
22.7805, 46.5539, 21.7024, 22.3656, 29.4675,
13.11403333, 13.7465, 39.9561, 43.7502, 34.3872,
33.2007, 32.85994444, 33.6577, 47.6333, 14.7271,
31.0306, 41.1486, 52.4563, 33.1355, 65.3387,
33.2787, 69.824, 21.7589, 52.153, 64.1194,
35.777, 41.6999, 39.4282, 29.6584, 39.3065,
14.8634, 33.0255, 17.00833333])
```

```
In [19]: rf.score(xtrain, ytrain)
```

```
Out[19]: 0.9850248244289491
```

```
In [20]: rf.score(xtest, ytest)
```

```
Out[20]: 0.8881757199749859
```

Gradient Boosting

```
In [21]: gbr = GradientBoostingRegressor()  
gbr.fit(xtrain, ytrain)
```

```
Out[21]: GradientBoostingRegressor()
```

```
In [22]: ypred3 = gbr.predict(xtest)
ypred3
```

```
Out[22]: array([48.79157614, 45.54661004, 70.55119948, 34.34700738, 12.57970481,
 40.47818786, 25.31268308, 50.76967494, 31.93648745, 42.08529343,
 38.8621663 , 16.88313428, 40.55450077, 42.63018298, 28.743147 ,
 22.08468511, 36.66677553, 19.71221872, 38.36618599, 32.22510605,
 39.06111309, 37.90260238, 47.62978022, 11.50798563, 36.89210508,
 34.54191527, 9.9434678 , 45.54085556, 53.39123654, 13.45318438,
 49.89658183, 35.97461115, 45.62408406, 56.81854747, 20.78863594,
 35.58599082, 31.93648745, 40.82795462, 12.32251451, 48.22571786,
 15.2637583 , 8.42836211, 36.87536899, 50.82452212, 13.25731205,
 75.04755107, 50.56520848, 34.83406847, 25.46069308, 9.35281935,
 48.23671782, 40.44595022, 25.16715084, 18.93063112, 41.32776619,
 34.83159316, 27.33408043, 9.9231908 , 36.95302139, 24.88275824,
 40.82795462, 14.97638632, 37.65193371, 50.3932738 , 30.33028221,
 25.47714514, 32.72309332, 17.08434351, 32.39431568, 23.90366913,
 11.60872804, 27.75774062, 7.68690474, 38.8621663 , 27.04821187,
 11.80772784, 49.58545467, 51.63631315, 52.50101904, 11.80951494,
 35.0102211 , 46.36125503, 38.1586267 , 39.42332254, 41.01707537,
 34.34700738, 34.40916587, 39.27484148, 28.99726436, 20.621892 ,
 33.22284073, 69.4375338 , 11.51456028, 51.34342972, 40.73389482,
 48.40862693, 25.69744285, 34.12751666, 24.35964127, 34.25965809,
 31.89998273, 40.48719032, 39.20450436, 22.22254031, 71.35150428,
 8.99596441, 53.01950472, 32.76702984, 36.0434034 , 56.97473765,
 35.0102211 , 48.89290123, 26.31788428, 40.04163257, 37.75710863,
 58.12238442, 16.74982854, 36.02954944, 57.3208685 , 40.73389482,
 23.33009288, 30.04593434, 48.93141617, 37.09424371, 20.9617292 ,
 45.01201301, 52.00060342, 49.14932167, 38.86484036, 37.31459056,
 10.51908811, 36.18071214, 22.6487033 , 71.15807006, 8.81247811,
 40.81732418, 31.43821331, 40.27705188, 22.13535075, 30.35546974,
 30.25390588, 34.64301274, 16.00248549, 34.84716865, 47.92073161,
 41.77221118, 28.2196103 , 16.80991801, 10.00591693, 15.82309007,
 41.48305422, 18.98844796, 39.67168448, 22.08878516, 43.51511577,
 38.19460066, 32.77658377, 72.75839664, 45.54661004, 42.60706283,
 14.5710747 , 55.80931861, 35.27998547, 50.05259892, 34.29400574,
 60.03853919, 43.33424919, 23.90366913, 42.74149456, 20.59711195,
 42.97624749, 14.06376191, 21.40688163, 38.11953379, 33.1698284 ,
 10.28383819, 30.51803067, 48.67907261, 42.01116553, 39.44324535,
 30.82544991, 32.5163854 , 56.97473765, 39.4758198 , 43.78591726,
 49.58545467, 26.17796899, 33.30497423, 29.34155182, 15.17357647,
 27.11126641, 40.55369244, 41.73052927, 47.12501063, 30.62866742,
 30.38130761, 53.72732516, 18.98844796, 20.48567335, 30.78840502,
 44.17603068, 58.12238442, 24.83752087, 44.34092253, 62.28484077,
 38.43134955, 28.94785385, 55.70418198, 24.78390486, 33.33371005,
 13.99723767, 27.93718963, 26.57295732, 45.06257072, 39.50827084,
 5.64639427, 16.80991801, 38.72772597, 23.19695729, 65.98106327,
 29.94705376, 50.43626666, 55.05803641, 22.63511197, 46.94824505,
 23.33009288, 45.6226325 , 23.73156935, 21.72480298, 29.38963876,
 12.24769018, 14.66360968, 40.70522513, 38.86484036, 35.85519276,
 34.83406847, 32.3345367 , 36.84770997, 44.78673872, 17.56991468,
 32.107081 , 39.40127577, 49.31515593, 31.0108153 , 67.36753659,
 33.1926849 , 71.10375148, 18.95160959, 50.15019853, 65.43266811,
 37.18100454, 36.83022521, 42.20293317, 25.3473084 , 36.83022521,
 16.75966806, 31.26796159, 17.90292473])
```



```
In [23]: gbr.score(xtrain, ytrain)
```

```
Out[23]: 0.9473233609145324
```

```
In [24]: gbr.score(xtest, ytest)
```

```
Out[24]: 0.8876649792898443
```

Evaluation Metrics

```
In [25]: models = pd.DataFrame(columns=["Model", "MAE", "MSE", "r2 Score", "RMSE"])
```

```
In [26]: mse_test1 = mean_squared_error(ytest, ypred1)
rmse_test1 = sqrt(mse_test1)
mae_test1 = mean_absolute_error(ytest, ypred1)
r2_test1 = r2_score(ytest, ypred1)
print("Mean Squared Error (MSE) of Linear Regression:", mse_test1)
print("Root Mean Squared Error (RMSE) of Linear Regression:", rmse_test1)
print("Mean Absolute Error (MAE) of Linear Regression:", mae_test1)
print("R-squared (R2) Score of Linear Regression:", r2_test1)

new_row = {"Model": "Linear Regression", "MAE": mae_test1, "MSE": mse_test1, "r2": r2_test1, "RMSE": rmse_test1}
models = models.append(new_row, ignore_index=True)
```

```
Mean Squared Error (MSE) of Linear Regression: 101.58139562951938
Root Mean Squared Error (RMSE) of Linear Regression: 10.07875962752954
Mean Absolute Error (MAE) of Linear Regression: 7.987048267733713
R-squared (R2) Score of Linear Regression: 0.6249829353885574
```

```
C:\Users\fahim\AppData\Local\Temp\ipykernel_8100\3682360161.py:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
  models = models.append(new_row, ignore_index=True)
```

```
In [27]: mse_test2 = mean_squared_error(ytest, ypred2)
rmse_test2 = sqrt(mse_test2)
mae_test2 = mean_absolute_error(ytest, ypred2)
r2_test2 = r2_score(ytest, ypred2)
print("Mean Squared Error (MSE) of Random Forest:", mse_test2)
print("Root Mean Squared Error (RMSE) of Random Forest:", rmse_test2)
print("Mean Absolute Error (MAE) of Random Forest:", mae_test2)
print("R-squared (R2) Score of Random Forest:", r2_test2)

new_row = {"Model": "Random Forest Regressor", "MAE": mae_test2, "MSE": mse_test2}
models = models.append(new_row, ignore_index=True)
```

Mean Squared Error (MSE) of Random Forest: 30.289998781726176
Root Mean Squared Error (RMSE) of Random Forest: 5.50363505164779
Mean Absolute Error (MAE) of Random Forest: 3.7622093567737203
R-squared (R2) Score of Random Forest: 0.8881757199749859

C:\Users\fahim\AppData\Local\Temp\ipykernel_8100\844649544.py:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
models = models.append(new_row, ignore_index=True)

```
In [28]: mse_test3 = mean_squared_error(ytest, ypred3)
rmse_test3 = sqrt(mse_test3)
mae_test3 = mean_absolute_error(ytest, ypred3)
r2_test3 = r2_score(ytest, ypred3)
print("Mean Squared Error (MSE) of Gradient Boosting:", mse_test3)
print("Root Mean Squared Error (RMSE) of Gradient Boosting:", rmse_test3)
print("Mean Absolute Error (MAE) of Gradient Boosting:", mae_test3)
print("R-squared (R2) Score of Gradient Boosting:", r2_test3)

new_row = {"Model": "Gradient Boosting Regressor", "MAE": mae_test3, "MSE": mse_test3}
models = models.append(new_row, ignore_index=True)
```

Mean Squared Error (MSE) of Gradient Boosting: 30.42834382385169
Root Mean Squared Error (RMSE) of Gradient Boosting: 5.516189248371713
Mean Absolute Error (MAE) of Gradient Boosting: 4.1151631863120315
R-squared (R2) Score of Gradient Boosting: 0.8876649792898443

C:\Users\fahim\AppData\Local\Temp\ipykernel_8100\602836561.py:11: FutureWarning: The frame.append method is deprecated and will be removed from pandas in a future version. Use pandas.concat instead.
models = models.append(new_row, ignore_index=True)

Hyperparameter Tuning with Grid Search (Random Forest)

```
In [29]: rfr_params = {  
    'n_estimators': [10, 50, 100, 200],  
    'max_depth': [3, 5, 7],  
    'min_samples_split': [2, 5, 10],  
    'min_samples_leaf': [1, 2, 4],  
    'max_features': ['auto', 'sqrt', 'log2'],  
    'bootstrap': [True, False],  
}
```

```
In [30]: rfr_grid_search = GridSearchCV(estimator = rf, param_grid= rfr_params, cv=5)  
rfr_grid_search.fit(xtrain, ytrain)
```


```
Out[30]: GridSearchCV(cv=5, estimator=RandomForestRegressor(),  
    param_grid={'bootstrap': [True, False], 'max_depth': [3, 5, 7],  
    'max_features': ['auto', 'sqrt', 'log2'],  
    'min_samples_leaf': [1, 2, 4],  
    'min_samples_split': [2, 5, 10],  
    'n_estimators': [10, 50, 100, 200]})
```

```
In [31]: rslt_tuning_rfr = pd.DataFrame(rfr_grid_search.cv_results_)
rslt_tuning_rfr
```

Out[31]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_bootstrap	param_max_
0	0.034372	0.006261	0.006254	0.007659	True	
1	0.141214	0.010873	0.009379	0.007658	True	
2	0.283549	0.004642	0.018748	0.006248	True	
3	0.578090	0.022100	0.034376	0.006248	True	
4	0.031248	0.000014	0.003125	0.006249	True	
...
643	0.481212	0.020725	0.037507	0.012494	False	
644	0.031248	0.000014	0.000000	0.000000	False	
645	0.121859	0.006256	0.009378	0.007657	False	
646	0.262480	0.020732	0.028126	0.006251	False	
647	0.480754	0.016287	0.040618	0.007667	False	

648 rows × 19 columns



```
In [32]: best_params = rfr_grid_search.best_params_  
best_params
```

```
Out[32]: {'bootstrap': True,  
          'max_depth': 7,  
          'max_features': 'auto',  
          'min_samples_leaf': 1,  
          'min_samples_split': 2,  
          'n_estimators': 100}
```

```
In [33]: best_model = rfr_grid_search.best_estimator_
```

```
In [34]: y_pred_rfr_gs = best_model.predict(xtest)
```

```
In [35]: y_pred_rfr_gs
```

```
Out[35]: array([47.92040099, 40.52689799, 71.58211124, 33.30713129, 11.58012633,
 42.46876169, 26.28671197, 47.73690835, 37.02743505, 40.90293881,
 42.86364698, 17.23928721, 39.03755841, 34.98342934, 24.35606468,
 20.07804879, 37.06631707, 19.11683594, 39.48535968, 31.12842398,
 38.54001272, 36.78046392, 46.36319482, 11.88506505, 33.86237226,
 38.56498511, 13.390136 , 49.40488611, 51.01214355, 14.24972005,
 62.60675558, 35.36346586, 43.61712052, 48.34306236, 17.15544563,
 40.74890933, 35.66424705, 42.17598621, 9.54397676, 49.78763691,
 16.25996869, 6.86998242, 40.70019778, 51.87030186, 13.29296988,
 65.33250554, 53.30765637, 34.52427976, 28.69068471, 8.90415592,
 52.19284842, 43.66099389, 27.21925191, 17.72016279, 46.39054872,
 33.82132875, 27.86853132, 13.79057243, 39.27360041, 22.2058718 ,
 42.17598621, 14.45612503, 35.5130445 , 50.19227302, 33.83554313,
 28.68416885, 35.61100011, 13.9236743 , 31.64764356, 24.05670501,
 13.50420175, 32.89872998, 11.55859256, 42.86364698, 27.24887848,
 12.21370646, 47.80951122, 48.17573163, 55.82313281, 10.20302867,
 36.25564988, 46.50105405, 38.21891958, 38.15472676, 42.46447782,
 33.30713129, 36.52153234, 36.67042981, 26.73199853, 23.12401719,
 35.16067855, 68.33573772, 13.51029729, 53.66460965, 40.9610554 ,
 46.61626102, 23.93984945, 35.38068868, 22.41267187, 33.99933317,
 31.22013704, 40.0611515 , 40.96761045, 23.8818081 , 65.09949935,
 13.64770294, 51.35885289, 33.21842673, 39.72551734, 59.48516733,
 36.43173968, 47.8942967 , 23.38042335, 41.40222476, 35.09883924,
 55.22830327, 17.70323862, 35.55646813, 55.45540609, 41.04596117,
 21.12562582, 30.60969678, 50.22869353, 33.38339385, 24.47508253,
 44.7856575 , 49.91049712, 48.50688425, 42.62526109, 36.8225774 ,
 13.79057243, 35.21468257, 23.30993524, 66.73499127, 12.31059627,
 42.37399549, 34.99782024, 37.28298175, 22.02683504, 33.32040566,
 30.42969016, 35.9596552 , 14.9481429 , 33.71856623, 45.46418058,
 41.217884 , 27.99483296, 15.01260838, 7.62266816, 17.37935515,
 41.49949564, 26.34482487, 39.68881287, 22.31613359, 42.46157879,
 36.26213173, 32.62133874, 68.37420028, 40.52689799, 45.99409747,
 12.65118812, 59.98990631, 34.28819198, 49.21855278, 34.62982587,
 62.06082154, 42.42920535, 23.68515493, 39.08164617, 21.62060427,
 43.21524764, 11.1354553 , 23.4335398 , 36.66280253, 34.02512821,
 12.9684464 , 33.71150347, 46.07194305, 46.5112073 , 36.57048803,
 30.37600642, 26.52502875, 59.48516733, 35.07243702, 40.03612085,
 47.76735122, 23.74757828, 38.07788974, 29.23310835, 15.62559241,
 25.96080667, 39.49812086, 40.87055162, 39.64862032, 30.99670422,
 29.23995774, 46.49671371, 26.34482487, 24.54703257, 32.06808215,
 39.46344585, 55.22830327, 26.93688172, 34.32689326, 63.71352667,
 36.11470824, 31.91645957, 60.24209629, 22.38402608, 35.50569784,
 15.34590963, 32.84494128, 23.77471552, 44.24840492, 42.82459225,
 6.38300376, 14.25210019, 36.65799178, 23.34381347, 71.79307975,
 29.01008937, 53.23498361, 59.00823232, 21.43656705, 51.04235683,
 22.06421068, 47.29078155, 25.51488152, 23.36946214, 28.57467227,
 13.63823042, 14.95830948, 39.36615103, 42.72550553, 34.28769567,
 34.52427976, 34.36583766, 33.40812857, 46.99882536, 14.93164109,
 30.90980457, 40.9864536 , 51.53565499, 29.35585065, 64.54482168,
 31.48348348, 66.89431028, 21.68709548, 49.14712962, 63.46157971,
 35.36754376, 40.54293824, 40.01476418, 32.40846245, 40.03231918,
 15.32957327, 33.83554313, 16.6569958 ])
```

```
In [36]: print("Accuracy of Random Forest after using Grid Search:")
rfr_grid_search.best_score_
```

Accuracy of Random Forest after using Grid Search:

Out[36]: 0.8793444216816084

```
In [37]: models.sort_values(by="RMSE", ascending= False)
```

Out[37]:

	Model	MAE	MSE	r2 Score	RMSE
0	Linear Regression	7.987048	101.581396	0.624983	10.078760
2	Gradient Boosting Regressor	4.115163	30.428344	0.887665	5.516189
1	Random Forest Regressor	3.762209	30.289999	0.888176	5.503635

Reasons behind on choosen Hyperparameters

I have applied hyperparameter on Random Forest Regression model. Here's a condensed explanation of the chosen hyperparameters and their reasoning for my concrete strength project:

- i) Number of trees (n_estimators): I am exploring a range of 10 to 200 trees to balance model complexity and accuracy. More trees often lead to better performance but also increase training time.
- ii) Tree depth (max_depth): Limiting maximum depth to 3, 5, or 7 to prevent overfitting, as excessively deep trees can memorize noise in the data instead of generalizing well.
- iii) Split requirements (min_samples_split, min_samples_leaf): You're setting minimum sample thresholds to control tree growth and reduce overfitting. Higher values make trees more conservative in splitting, ensuring patterns are based on sufficient data.
- iv) Feature sampling (max_features): You're experimenting with different strategies to diversify trees and prevent reliance on a few dominant features. 'auto' uses a default heuristic, while 'sqrt' and 'log2' limit features considered at each split.
- v) Bootstrapping (bootstrap): I'm testing both with and without bootstrapping to assess its impact on accuracy and diversity. Bootstrapping involves sampling data with replacement for each tree, potentially improving robustness but also increasing randomness.

Essentially, I'm exploring a variety of hyperparameter settings to fine-tune the model's balance between complexity, accuracy, and generalization to achieve optimal performance in predicting concrete strength.

Comparative Analysis

Here highest accuracy gain by using Random Forest Regression model. The accuracy is 89%.

The limitations of each model:-

Limitations of Linear Regression:

- i) Assumes linearity, might not capture complex relationships or interactions between features.
- ii) Sensitive to outliers, which can significantly affect predictions.

Limitations of Random Forest:

- i) Less interpretable than linear regression, harder to understand feature importance.
- ii) Can be computationally expensive to train, especially with large datasets.

Limitations of Gradient Boosting:

- i) More prone to overfitting than Random Forest if not tuned properly.
- ii) Computationally expensive to train, especially with large datasets.
- iii) Less interpretable than linear regression.

Conclusion

In conclusion, the concrete strength prediction project has demonstrated significant advancements in the field of construction materials engineering. Through the utilization of advanced machine learning algorithms and extensive datasets, we have successfully developed a robust model capable of accurately predicting concrete strength.

The accuracy and reliability of our machine learning model were validated through rigorous testing and comparison with traditional methods of concrete strength prediction. The model not only outperformed existing approaches but also showcased its adaptability to diverse scenarios and variations in material composition.

In essence, the concrete strength prediction machine learning project stands as a beacon of innovation, offering a glimpse into the future of construction materials engineering and paving the way for more sophisticated and precise methodologies in the field.